

Blog Profiles

Reference for the scoring profiles used by [horizon-blog](#). Each profile defines the audience, gate path logic, scoring dimensions, and score milestones that determine which items are selected and ranked.

Update this file when profile code in [src/blog/profiles/](#) changes.

News Profile

Audience: General tech readers — people who follow tech news but are not necessarily practitioners.
Purpose: Selects and ranks items for their broad significance, newsworthiness, and accessibility to a non-expert audience.

Gate path

A single path ([news_stories](#)); items must clear all three dimensions:

[news_stories](#): significance ≥ 6 AND newsworthiness ≥ 5 AND narrative_clarity ≥ 4

Scoring dimensions & weights

Dimension	Description	Gate	Weight
significance	How broadly significant is this development for the AI/tech landscape?	≥ 6	0.45
newsworthiness	Is this timely, original, and a primary announcement rather than derivative commentary?	≥ 5	0.35
narrative_clarity	Is there a clear, compelling story a non-expert reader can follow?	≥ 4	0.20

Weighted sum: $0.45 \times \text{significance} + 0.35 \times \text{newsworthiness} + 0.20 \times \text{narrative_clarity}$

Score milestones

[significance](#) — gate ≥ 6

Score	Meaning
1	Niche announcement affecting a tiny audience
5	Noteworthy for a specific community but limited broader impact
6 (gate)	Broadly relevant — affects AI/tech users or the industry in a tangible way
8	Affects many users, products, or industries in a meaningful way
10	Industry-defining event — affects everyone in the tech ecosystem

[newsworthiness](#) — gate ≥ 5

Score	Meaning
1	Old news resurfacing or pure opinion with no new facts
5 (gate)	Secondary coverage of a genuine event with some new angle
8	Primary announcement, breaking news, or original reporting
10	Major breaking development from the primary source, first of its kind

[narrative_clarity](#) — gate ≥ 4 (*rarely binding in practice*)

Score	Meaning
1	Highly technical or jargon-heavy with no accessible angle
5	Story exists but requires significant background to appreciate
8	Clear narrative with obvious stakes a general reader would understand
10	Compelling human or societal story that writes itself for any audience

Engineer Profile

Audience: ML/MLOps engineers — people who build, fine-tune, serve, and evaluate models in production.
Purpose: Selects items based on ML engineering relevance and technical substance. Two named gate paths let research-significant posts through alongside immediately deployable content.

Gate paths

An item passes if it clears **all dimensions in any one path** (OR across paths, AND within a path):

Path	Purpose	Gate conditions
engineering_applicability	Ready-to-adopt tools, APIs, model releases, and production guides	ml_engineering_relevance ≥ 7 AND technical_substance ≥ 5 AND production_applicability ≥ 6 AND ai_ecosystem_significance ≥ 7
technical_insights	Research with important engineering insights, even if not yet deployable	ml_engineering_relevance ≥ 8 AND technical_substance ≥ 5 AND engineering_insight ≥ 6

When an item passes both paths, **engineering_applicability** is used (first match) and its weighted sum is applied.

Scoring dimensions & weights

Dimension	Description	engineering_applicability gate	engineering_applicability weight	technical_insights gate	technical_insights weight
ml_engineering_relevance	Relevance to building, fine-tuning, serving, or evaluating ML models in production	≥ 7	0.35	≥ 8	0.45
technical_substance	Concrete technical detail an ML engineer can learn from	≥ 5	0.30	≥ 5	0.30
production_applicability	Can be adopted in an ML stack this sprint	≥ 6	0.20	—	—
ai_ecosystem_significance	Major model/API release from a key provider vs. niche	≥ 7	0.15	—	—
engineering_insight	Advances understanding of how to design, build, or operate ML systems	—	—	≥ 6	0.25

Weighted sums:

- engineering_applicability: $0.35 \times \text{ml_engineering_relevance} + 0.30 \times \text{technical_substance} + 0.20 \times \text{production_applicability} + 0.15 \times \text{ai_ecosystem_significance}$
- technical_insights: $0.45 \times \text{ml_engineering_relevance} + 0.30 \times \text{technical_substance} + 0.25 \times \text{engineering_insight}$

Score milestones

ml_engineering_relevance — gate ≥ 7 (engineering_applicability) / ≥ 8 (technical_insights)

Score	Meaning
1	No ML engineering content — business, policy, or consumer product news
5	Tangentially ML-related (platform news, infra update with minor ML angle)
7 (engineering_applicability gate)	Relevant to ML engineering — about a tool, technique, or result engineers work with
8 (technical_insights gate)	Directly about a technique, tool, or result engineers should act on
10	Paradigm shift in how models are trained, served, or evaluated

technical_substance — gate ≥ 5 (both paths)

Score	Meaning
1	Pure business/PR with no technical details — partnership deal, funding round, executive hire
4	Vague capability claims with no supporting numbers, architecture details, or methodology; OR feature announced with no technical specifics and no confirmed availability
5 (gate — both paths)	Research or lab blog post that describes a method or finding at a conceptual level — some methodology described, but key implementation details or reproducible numbers are not provided
6	Technically rich content an ML engineer can learn from: new open or API-accessible model with architecture details, efficiency numbers, or benchmark results; OR working SDK/API with technical docs; OR engineering blog with concrete implementation details; OR research algorithm with experimental results
7	Deployed model with model card and benchmark results; OR paper with clear methodology, concrete experiments, and reproducible numbers
9	Open-weights model with full technical report and benchmark suite, OR paper + code + benchmark methodology
10	Full paper + open-source code + benchmark + ablations + model weights

production_applicability — gate ≥ 6 (engineering_applicability only)

Score	Meaning
1	Theoretical or hypothetical — years from being usable
5	Limited/beta access or requires significant custom work
6 (gate)	Accessible with moderate effort — integration path exists but needs custom work
8	Available with straightforward setup, clear integration path
10	Available now, works out of the box, engineers can adopt immediately

ai_ecosystem_significance — gate ≥ 7 (engineering_applicability only)

Score	Meaning
1	Internal tooling or niche product with limited audience
5	Minor model variant or incremental version bump
7 (gate)	Meaningful release from a key provider that practitioners would evaluate for deployment
8	Primary announcement of a significant model update or major new capability from a key provider
10	Flagship model release (GPT-5, Claude 4, Llama 4, Gemini 2) or paradigm-shifting product change

engineering_insight — gate ≥ 6 (technical_insights only)

Score	Meaning
1	No engineering insight — pure business news, market commentary, or consumer product update
4	Broadly interesting to ML people but does not advance engineering knowledge
6 (gate)	Directly informs how practitioners should think about designing, evaluating, or operating ML systems
8	Hands-on research with clear engineering implications practitioners should incorporate into their workflow
10	Paradigm-shifting engineering insight with validated results that changes how ML systems should be built